

Chapter 7 (7.1 – 7.4)

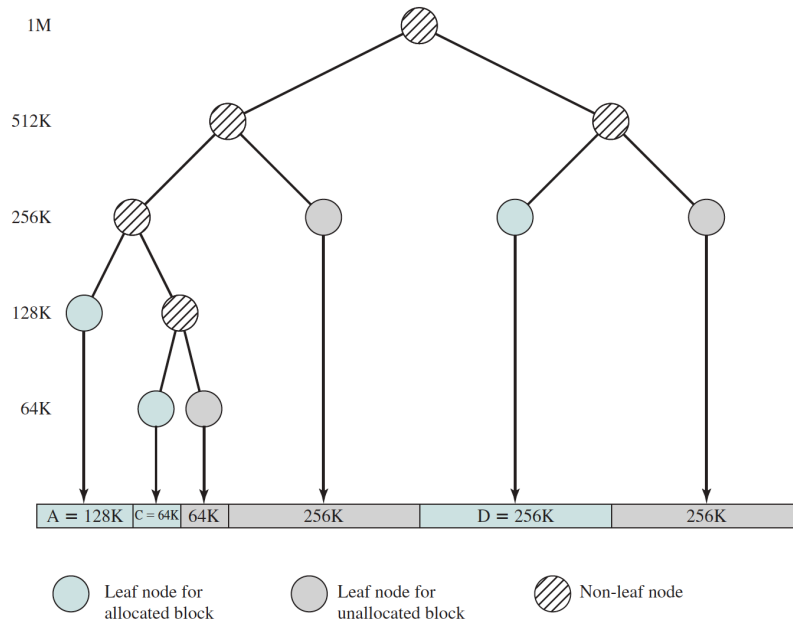
- **Frame** – fixed-length block of main memory
- **Page** – fixed-length block of data in secondary memory that can be copied into a frame
- **Segment** – variable-length block of data in secondary memory that can be copied into main memory, or split into pages then copied
- **Memory management** satisfies these requirements:
 - Relocation
 - Process may be swapped out and into a different memory location
 - Translation of memory references in program code
 - Protection
 - Process should not access memory in another process
 - Sharing
 - Allow processes to access same portion of memory
 - Multiple processes in one program need to access the same code
 - Logical organization
 - Modularized programs; different degrees of protection for each module
 - Physical organization
 - Overlaying allows modules to be assigned the same region of memory

Technique	Description	Strengths	Weaknesses
Fixed Partitioning	Main memory is divided into a number of static partitions at system generation time. A process may be loaded into a partition of equal or greater size.	Simple to implement; little operating system overhead.	Inefficient use of memory due to internal fragmentation; maximum number of active processes is fixed.
Dynamic Partitioning	Partitions are created dynamically, so that each process is loaded into a partition of exactly the same size as that process.	No internal fragmentation; more efficient use of main memory.	Inefficient use of processor due to the need for compaction to counter external fragmentation.

- **Fixed partitioning**
 - Placement in fixed-size partitions: pick any available
 - Placement in variable-size partitions: pick smallest-fitting partition
 - One process queue per partition vs. single queue for all partitions
 - **Internal fragmentation** – wasted memory within partitions
- **Dynamic partitioning**
 - Placement:
 - Best-fit – closest (and greater than) required size
 - First-fit – first available block from beginning
 - Next-fit – first available block from last placement location
 - **External fragmentation** – wasted memory between partitions

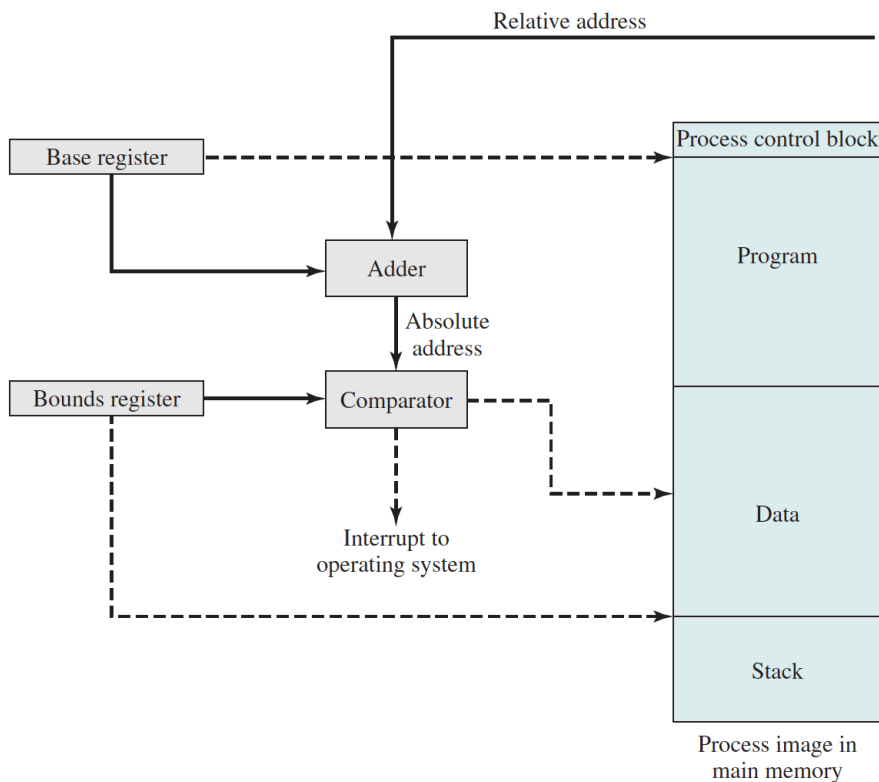
- Buddy system:

- Split memory in half until $B/2 \leq \text{required size} \leq B$, where B = block size
- When a block (size B)'s buddy is freed, coalesce into a block of size $2B$



- Relocation in partition:

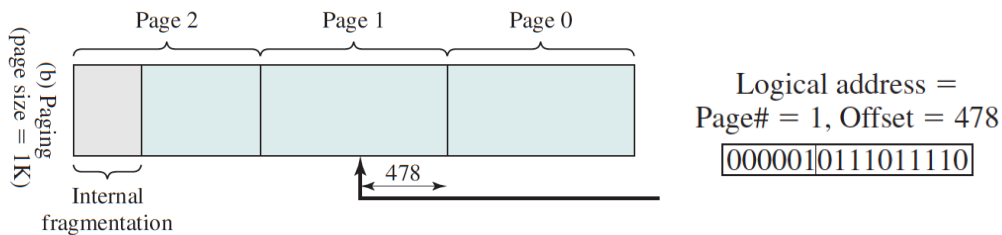
- Logical address = location relative to the beginning of program
- Relative address is added to base register to produce absolute address, and compared with bounds register before accessing data



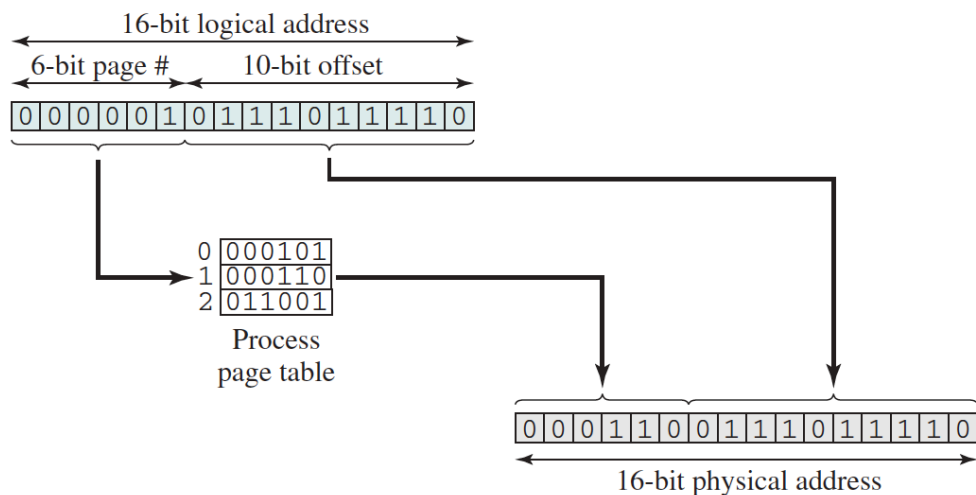
Simple Paging	Main memory is divided into a number of equal-size frames. Each process is divided into a number of equal-size pages of the same length as frames. A process is loaded by loading all of its pages into available, not necessarily contiguous, frames.	No external fragmentation.	A small amount of internal fragmentation.
Simple Segmentation	Each process is divided into a number of segments. A process is loaded by loading all of its segments into dynamic partitions that need not be contiguous.	No internal fragmentation; improved memory utilization and reduced overhead compared to dynamic partitioning.	External fragmentation.

- **Paging**

- Page table[page index] = frame index
- Address is split into n (page number) + m (offset) bits

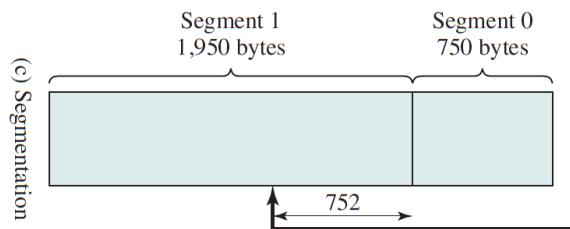


- Address translation:



- **Segmentation**

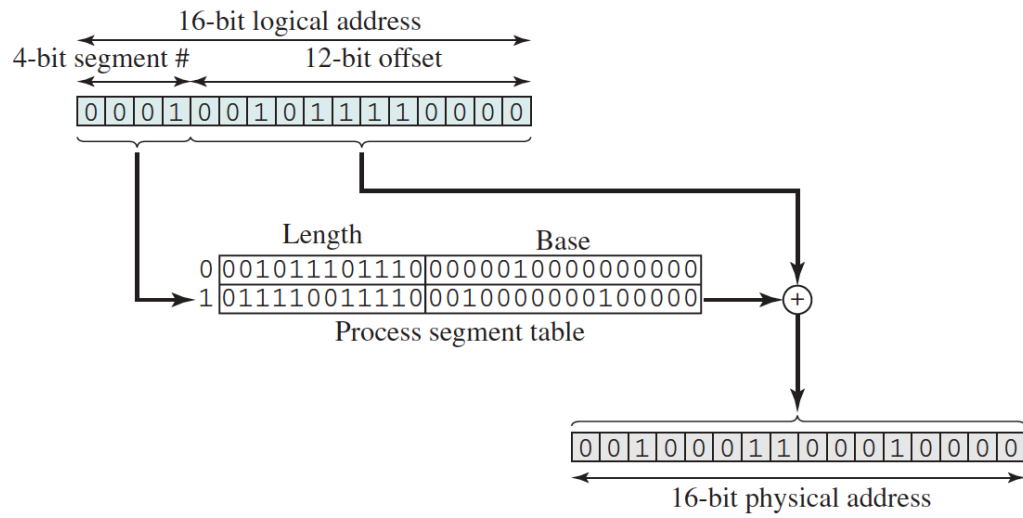
- Segment table[segment index] = length, base



Logical address =
Segment# = 1, Offset = 752

0001001011110000

- Address translation:



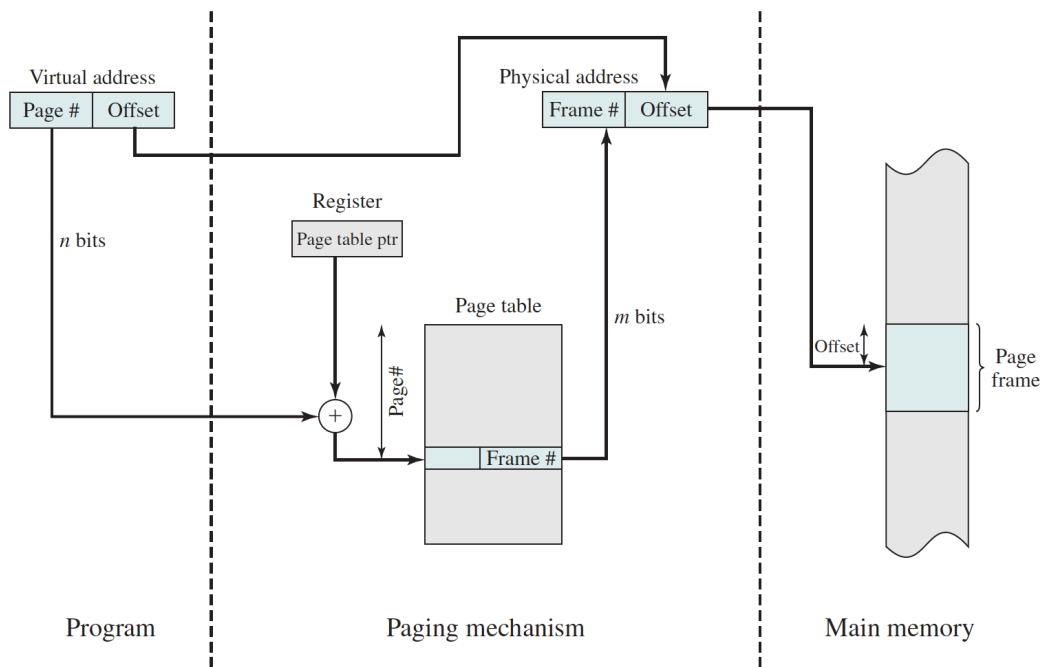
Chapter 8 (8.1 – 8.4)

- **Virtual memory** – storage allocation scheme in which secondary memory can be addressed as though it were part of main memory
 - Key benefits:
 - More processes can be maintained in main memory
 - A process may be larger than main memory
 - Resident set – portion of a process actually in main memory
 - When a mem reference outside of resident set is encountered (page fault), interrupt & block the process, issue I/O to retrieve the needed program
 - **Lazy loading** – defer loading until needed – keeps CPU busy, less wasted time
 - Virtual address – location in virtual memory that can be accessed as though it's in main memory
 - Real address – actual location in main memory

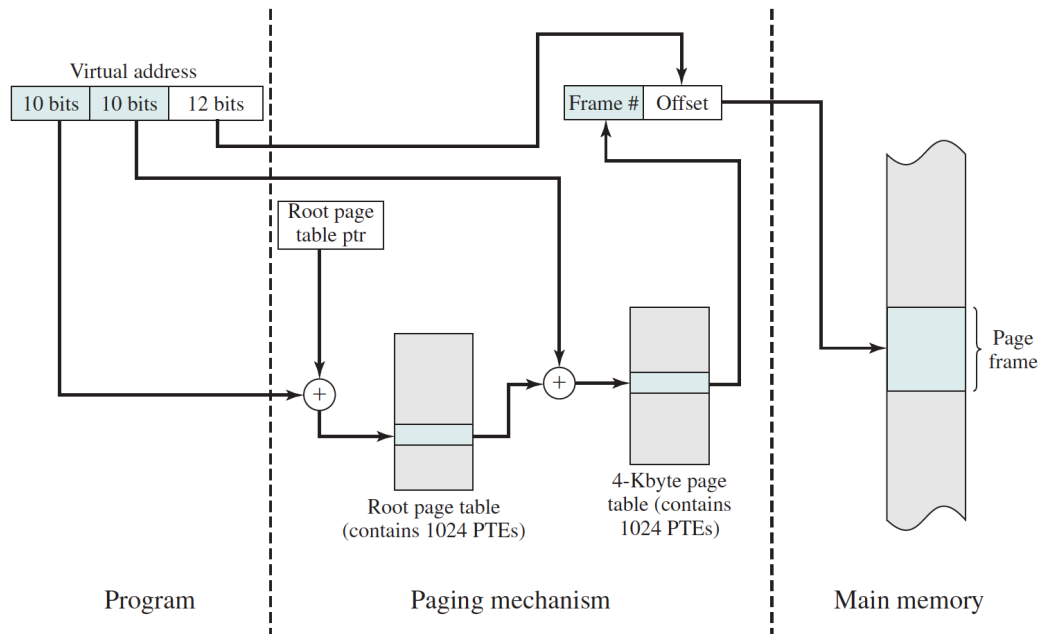
Virtual Memory Paging	As with simple paging, except that it is not necessary to load all of the pages of a process. Nonresident pages that are needed are brought in later automatically.	No external fragmentation; higher degree of multiprogramming; large virtual address space.	Overhead of complex memory management.
Virtual Memory Segmentation	As with simple segmentation, except that it is not necessary to load all of the segments of a process. Nonresident segments that are needed are brought in later automatically.	No internal fragmentation, higher degree of multiprogramming; large virtual address space; protection and sharing support.	Overhead of complex memory management.

- **Thrashing** – swap out something that is needed again immediately; wastes too much time on swapping
 - Solve with principle of locality

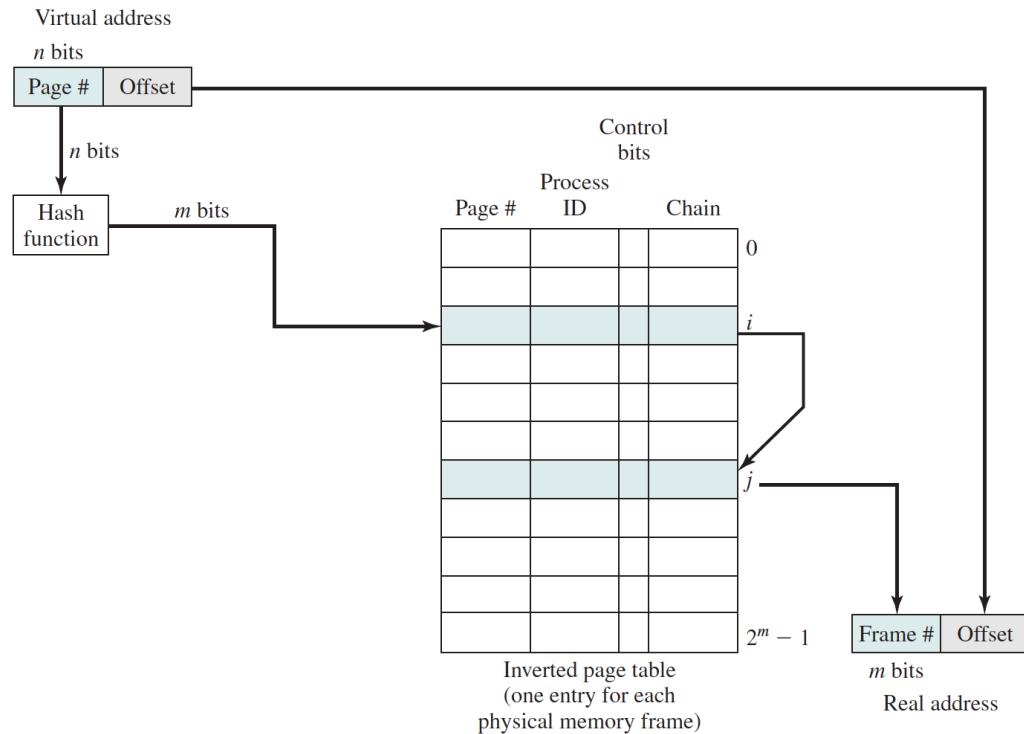
- Hardware support for virtual memory:
- **Page tables**
 - Entry contains present (P) & modified (M) control bits, in addition to frame #
 - Page tables are in virtual memory as well
 - Large page tables referenced by page directories
 - Page directories are in main memory
 - A “root table” for the pages that the actual page table occupies
- Single-level page table:



- Two-level page table:



- Inverted page table – one entry per real memory frame
 - Page table doesn't have to be paged since # of pages = # of frames
 - Page # is hashed into table index; overflows are chained

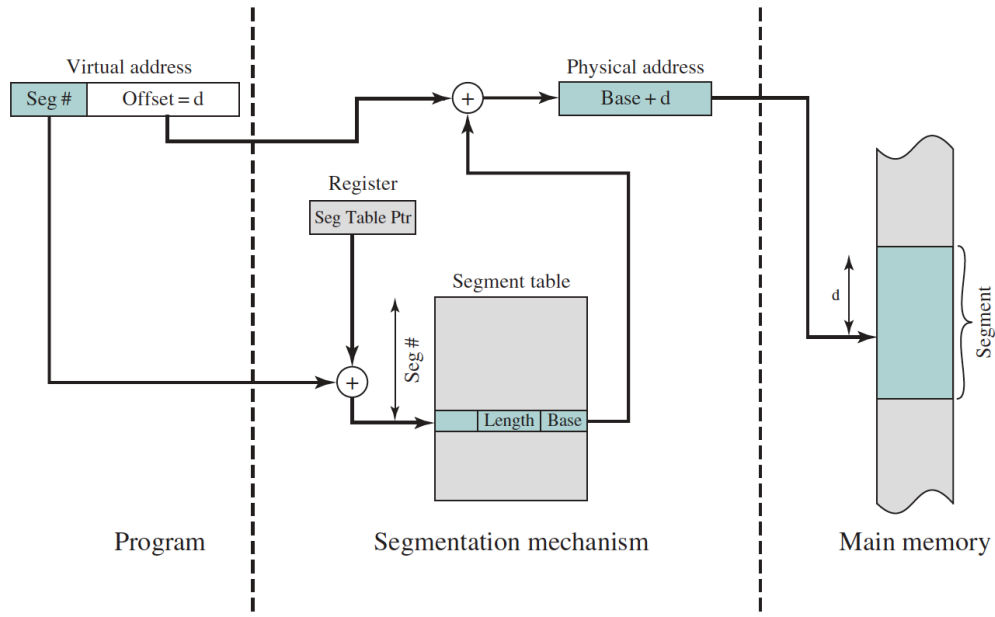


- Translation look-aside buffer – cache for page table entries
 - Check TLB →
 - If TLB hit → get frame # from TLB page entry
 - If TLB miss → lookup in page table
 - If present bit set → get frame # from page entry
 - If present bit not set → page fault
- Cache is also consulted when retrieving a block from memory
- Page size:
 - Large = more internal fragmentation; weaker principle of locality → more page faults
 - Small = larger page tables → more pages in virtual memory → more page faults
- Pages in memory:
 - More pages in memory → less page faults

- **Segmentation:**

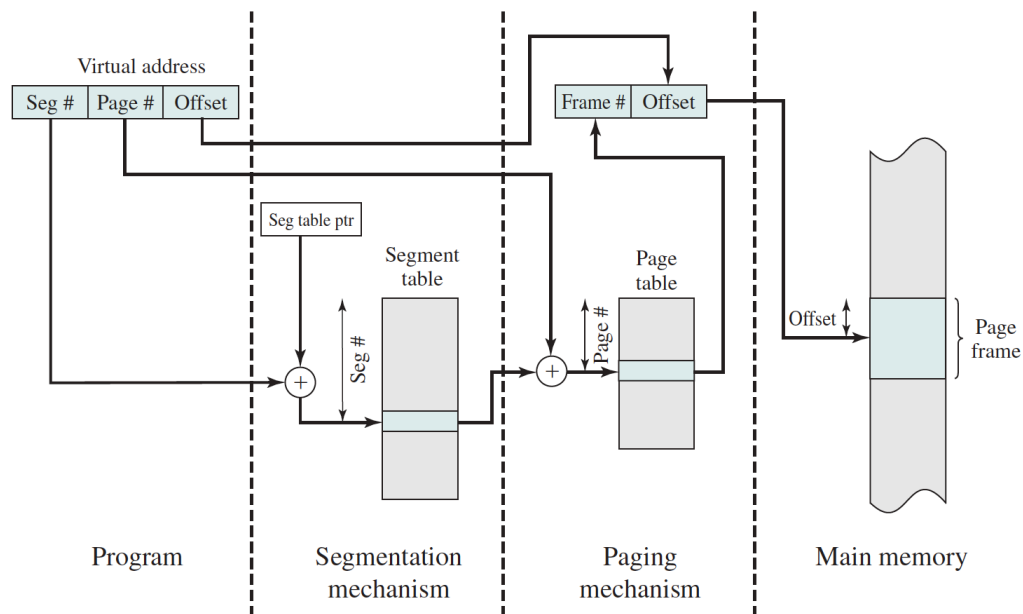
- Advantages:

- Simplifies growing data structures
 - Programs can be altered & recompiled independently
 - Lends itself to sharing
 - Lends itself to protection



- **Combined paging & segmentation:**

- Each segment is paged
 - Segment offset = page # & offset



- Memory management software:
- OS virtual memory policies:
 - **Fetch policy** – when to load pages
 - Demand paging – only load page when reference to it is made
 - Prepaging – load multiple contiguous pages
 - Good for processes that use many contiguous pages
 - Placement policy
 - Important in segmentation system – best-fit, first-fit, next-fit etc.
 - Not as relevant in paging/combined systems
 - **Replacement policy** – which page(s) in main memory to replace
 - Locked frames cannot be replaced (i.e. kernel frames)
 - Optimal (impossible in practice)
 - LRU
 - FIFO
 - Clock – look for page with use bit = 0
 - When searching, every use bit = 1 encountered is set to 0
 - Referencing a page sets use bit to 1
 - Clock policy with use bit & modified bit
 - 1. Search for $u = 0$ & $m = 0$
 - 2. Search for $u = 1$ & $m = 0$; set u to 0 when scanning
 - Repeat 1 and 2 if necessary
 - Page buffering
 - Replaced pages are not removed from memory; instead are added to free (unchanged) & modified lists
 - Lists act as caches, can be added back to resident set if referenced
 - Modified list is written out in clusters
 - **Resident set management**

	Local Replacement	Global Replacement
Fixed Allocation	<ul style="list-style-type: none"> • Number of frames allocated to a process is fixed. • Page to be replaced is chosen from among the frames allocated to that process. 	<ul style="list-style-type: none"> • Not possible.
Variable Allocation	<ul style="list-style-type: none"> • The number of frames allocated to a process may be changed from time to time to maintain the working set of the process. • Page to be replaced is chosen from among the frames allocated to that process. 	<ul style="list-style-type: none"> • Page to be replaced is chosen from all available frames in main memory; this causes the size of the resident set of processes to vary.

- Size:
 - Fixed allocation
 - Variable allocation
 - Processes w/ high page fault rate can be allocated more frames
- Scope:

- Local replacement – only choose from resident pages of faulting process
 - Global replacement – choose from any unlocked pages
- **Working set** – set of pages referenced by a process in the last given amount of time
 - Working set strategy:
 - Monitor the W of each process
 - Periodically remove pages in resident set but not in W (LRU)
 - Process may only execute if its W is in memory
 - Algorithms that follow this strategy:
 - Page fault frequency
 - If time since last page fault < threshold → grow resident set
 - If time since last page fault > threshold → remove all unused pages
 - Variable-interval sampled working set
- **Cleaning policy** – when to write out modified pages
 - Demand cleaning – written out when it's replaced
 - Precleaning – written out in batches
 - Page buffering – uses modified & unmodified lists of pages
- **Load control** – how many processes to maintain in memory (aka. multiprogramming level)
 - Too few → processes too often blocked
 - Too many → resident sets too small → thrashing (too much swapping)
- Process suspension – which process to take out of memory
 - Lowest-priority
 - Faulting process
 - Last activated
 - Smallest resident set
 - Largest process
 - Largest remaining execution window

- **UNIX/Solaris**

- Page replacement uses a refined, two-handed clock algorithm

- **Linux**

- 3-level page table structure
 - Page directory (in main memory) → middle directory → page table → frame #
 - Page replacement uses split LRU algorithm (active & inactive lists)