

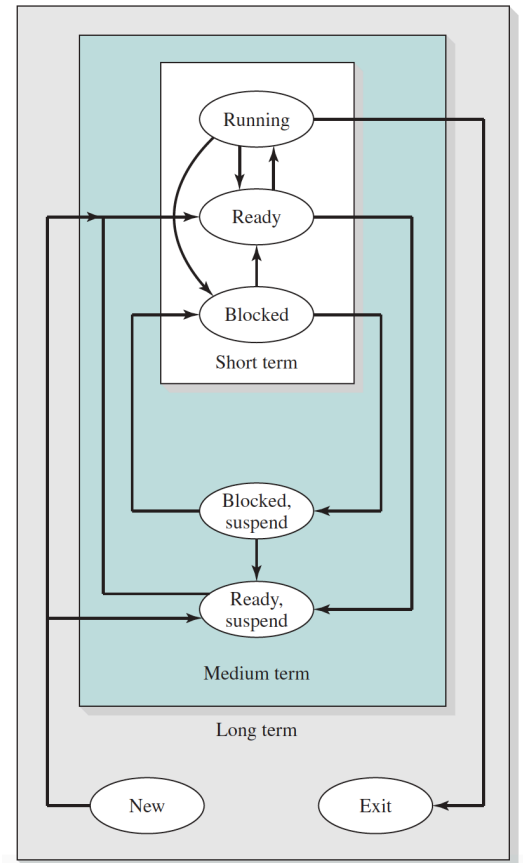
Chapter 9 (9.1 – 9.3)

- **Long-term scheduling** – deals with new processes
 - Can another process be created?
 - Which process to admit?
- **Medium-term scheduling** – deals with swapping
- **Short-term scheduling** – deals with which process to execute
 - Aka. the dispatcher
- **I/O scheduling** – which process's pending I/O request to handle

- Short-term scheduling criteria:

| | <i>User-oriented</i> | <i>System-oriented</i> |
|----------------------------|---|--|
| <i>Performance related</i> | <ul style="list-style-type: none">• Turnaround time• Response time• Deadlines | <ul style="list-style-type: none">• Throughput• Processor utilization |
| <i>Other</i> | <ul style="list-style-type: none">• Predictability | <ul style="list-style-type: none">• Fairness (avoid starvation)• Enforcing priorities• Balancing resources |

- Scheduling policies
 - Pure priority scheduling may cause starvation
 - Quantities:
 - w = time spent waiting so far
 - e = time spent executing so far
 - s = total service time (including e)
 - Turnaround time = waiting time + service time
 - Decision modes:
 - Non-preemptive – process will continue execution until it terminates or is blocked
 - Preemptive – currently running process may be pre-empted
 - First-come first-served
 - Bad for short processes
 - Round robin
 - Inefficient for I/O processes
 - Virtual round robin adds an auxiliary queue – dispatcher favours I/O processes that just became unblocked
 - Shortest process next – needs to know s
 - Long processes have less predictability, may be starved
 - Shortest remaining time – needs to know s
 - Long processes may still be starved
 - Better turnaround time than SPN
 - Highest response ratio next – needs to know s



- Picks based on normalized turnaround time
- Feedback
 - Every time a process is pre-empted it is demoted in priority
 - Long processes can be starved

| | FCFS | Round Robin | SPN | SRT | HRRN | Feedback |
|----------------------------|---|---|---|-----------------------------|------------------------------------|-------------------------------|
| Selection Function | $\max[w]$ | constant | $\min[s]$ | $\min[s - e]$ | $\max\left(\frac{w + s}{s}\right)$ | (see text) |
| Decision Mode | Non-preemptive | Preemptive (at time quantum) | Non-preemptive | Preemptive (at arrival) | Non-preemptive | Preemptive (at time quantum) |
| Throughput | Not emphasized | May be low if quantum is too small | High | High | High | Not emphasized |
| Response Time | May be high, especially if there is a large variance in process execution times | Provides good response time for short processes | Provides good response time for short processes | Provides good response time | Provides good response time | Not emphasized |
| Overhead | Minimum | Minimum | Can be high | Can be high | Can be high | Can be high |
| Effect on Processes | Penalizes short processes; penalizes I/O bound processes | Fair treatment | Penalizes long processes | Penalizes long processes | Good balance | May favor I/O bound processes |
| Starvation | No | No | Possible | Possible | No | Possible |

- Fair share scheduling
 - In some systems the process pool need to be regarded as a collection of process sets, each with some weighting that determines their share of system resources
- **Unix scheduling**
 - Uses multi-level feedback
 - Round robin within priority queues
 - Processes are pre-empted & priorities recalculated every second
 - Processes are restricted within their priority bands: (decreasing priority)
 - Swapper
 - Block I/O device control
 - File manipulation
 - Character I/O device control
 - User processes

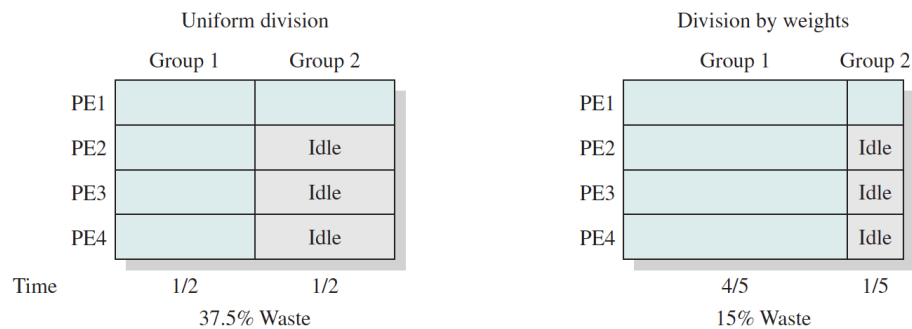
Chapter 10 (10.1 – 10.5)

- Types of multiprocessor systems
 - Loosely coupled/distributed multiprocessor (aka. cluster)
 - Functionally specialized processors
 - Tightly coupled multiprocessor
- **Granularity** – frequency of synchronization

| Grain Size | Description | Synchronization Interval (Instructions) |
|-------------------|--|--|
| Fine | Parallelism inherent in a single instruction stream | < 20 |
| Medium | Parallel processing or multitasking within a single application | 20–200 |
| Coarse | Multiprocessing of concurrent processes in a multiprogramming environment | 200–2,000 |
| Very Coarse | Distributed processing across network nodes to form a single computing environment | 2,000–1M |
| Independent | Multiple unrelated processes | Not applicable |

- Independent parallelism – no synchronization among processes
 - E.g. time sharing system
 - Coarse/very coarse – handled as set of concurrent processors on a uniprocessor
 - Can be supported on multiprocessor with little change
 - Medium/fine – coordination & interaction between threads in a program
- Design issues
 - Assignment of processes to processors
 - Static assignment vs. global queue vs. dynamic load balancing
 - Master/slave – kernel functions occupy a single processor
 - Master processor controls all memory & I/O, handles scheduling
 - Disadvantage: master can be a bottleneck; failure of master → failure of system
 - Peer architecture – kernel can execute on any processor; processor self-schedules
 - Disadvantage: complicates OS
 - Multiprogramming on individual processors
 - With many processors and higher granularity, it's no longer necessary for every processor to be busy all the time
 - Process dispatching
 - Sophisticated scheduling algorithms may be unnecessary
- **Thread scheduling**
 - Load sharing
 - Global queue of ready threads
 - Advantages:
 - No processor is left idle

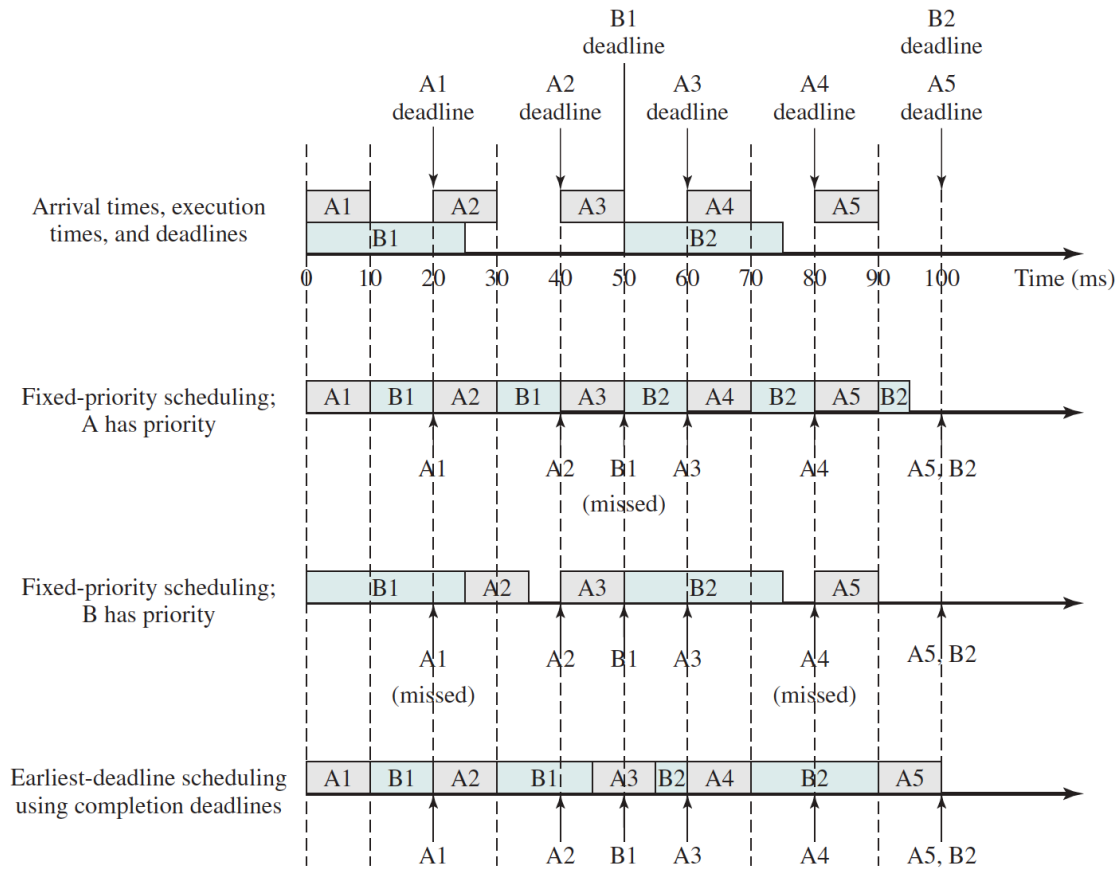
- No centralized scheduler required
- Disadvantages:
 - Central queue requires mutual exclusion – bottleneck
 - Preemptive threads tend to resume on another process – inefficient for caches
 - Unlikely all threads of a process will be active at the same time
- Gang scheduling
 - Set of threads simultaneously scheduled to set of processors
 - Great for highly synchronized threads (medium/fine grained)
 - Single-threaded applications can cause inefficient processor use



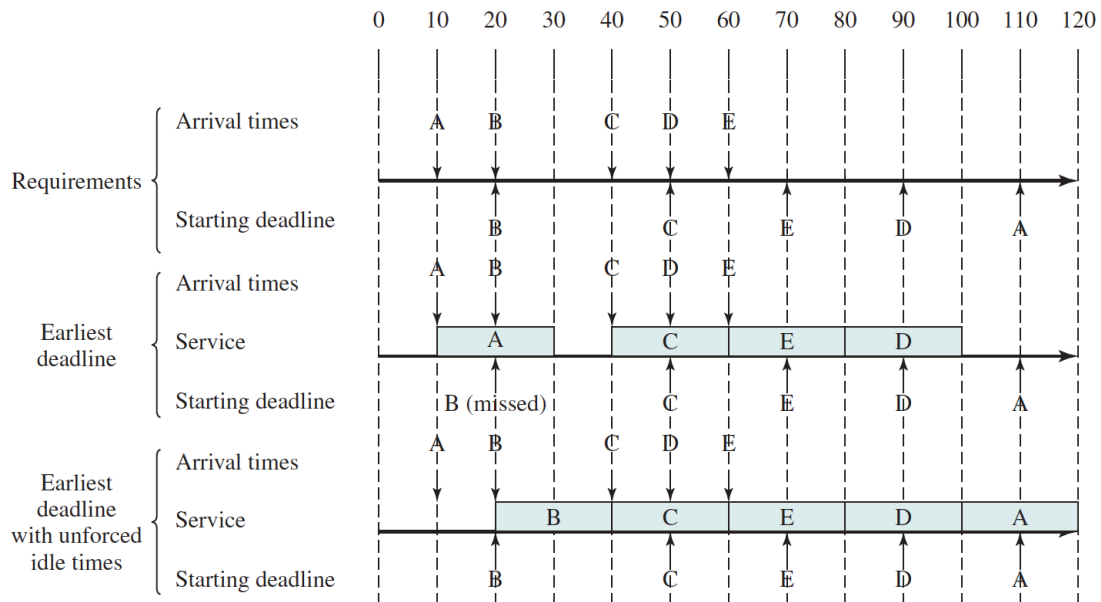
- Dedicated processor assignment
 - Dedicate group of processors to an application; each thread gets a processor
 - Useful for highly parallel, many processor system; no processor switching ever
 - Dynamic scheduling
 - # of threads in each application can be altered dynamically
 - In multicore systems, minimizing access to off-chip memory (i.e. caches) takes precedence over maximizing processor utilization
- **Real-time systems**
 - Correctness of the system depends on logical result as well as the time at which result is produced
 - Real-time tasks must be able to keep up with deadlines
 - Hard (necessary, else fault) vs. soft (desirable)
 - Periodic vs. aperiodic
- Characteristics of RT OS:
 - Determinism – perform operations at fixed, predetermined times/time intervals
 - Depends on interrupt response speed & capacity to handle requests
 - How long before interrupt is acknowledged
 - Responsiveness – after acknowledging, how long before interrupt is serviced
 - Depends on time to setup & perform the interrupt service routine, and interrupt nesting
 - User control – users have more fine-tuned control over priority, time, memory (e.g. paging), I/O, etc.
 - Reliability – degradation of performance have much more severe consequences
 - Fail-soft operation – ability to fail in a way that preserves as much capability and data as possible

- Stability – when unable to meet all deadlines, prioritize most critical tasks
- Common features of RT OS:
 - Stricter use of priorities
 - Interrupt latency is bounded & short
 - More precise and predictable timing
 - Fast process/thread switching
 - Preemptive priority-based scheduling
- **Real-time scheduling**
 - Static table-driven
 - Perform static analysis, produces a schedule that fits requirements
 - Useful for periodic tasks
 - Static priority-driven preemptive
 - Perform static analysis and assign priorities to tasks; let traditional scheduler schedule them
 - Dynamic planning-based
 - Feasibility of task & schedule determined at run-time
 - Task is only accepted if deadline is feasible
 - Dynamic best effort
 - No feasibility analysis; system tries to meet all deadlines at runtime
 - Commonly used in practice
- **Deadline scheduling**
 - Real-time is less concerned with speed, and more with completing tasks at the right time
 - Schedule tasks based on info about them:
 - Ready time
 - Starting deadline
 - Completion deadline
 - Processing time
 - Resource requirements
 - Priority
 - Subtask structure
 - Scheduling based on earliest (starting/completion) deadlines minimizes tasks that miss their deadlines (EDF)
 - Starting deadlines → use non-preemptive scheduling
 - Completion deadlines → use preemptive scheduling
 - Can achieve 100% CPU utilization with preemption

- Example:
 - A has execution time = 10, deadline every 20
 - B has execution time = 25, deadline every 50



- Example with aperiodic tasks:



- **Rate monotonic scheduling**

- Tasks with shorter periods → higher priority
- Commonly adopted in industry

- Priority inversion

- Circumstances in system forces a higher-priority task to wait for a lower-priority task
- E.g. lower-p locks a resource, and higher-p tries to lock the same resource
- Unbounded priority inversion
 - Duration of priority inversion depends on unpredictable actions of unrelated tasks
- Priority inheritance
 - Lower-p task temporarily inherits the priority of any higher-p task that is waiting on a resource they have
- Priority ceiling – resource & process that accesses it is given a temporary higher priority

- **Linux scheduling**

- SCHED_FIFO – FIFO real-time threads
- SCHED_RR – round robin real-time threads
 - Real-time threads have priorities 0 – 99
- SCHED_OTHER – non-real-time threads
 - Priorities 100 – 139
 - No preemption
- 140-bit priority array each for the active and expired priority queues (140 queues each)
 - Each array cell points to the queue for that priority
 - Aka. the O(1) scheduler

- **UNIX SVR4 scheduling**

- Queue for each priority is executed in round robin

| Priority class | Global value | Scheduling sequence |
|----------------|--------------|---------------------|
| Real time | 159 | First ↓ Last |
| | • | |
| | • | |
| | • | |
| Kernel | 100 | |
| | • | |
| | • | |
| | • | |
| Time shared | 99 | |
| | • | |
| | • | |
| | • | |
| | 60 | |
| | • | |
| | • | |
| | • | |
| | 59 | |
| | • | |
| | • | |
| | • | |
| | 0 | |
| | • | |
| | • | |
| | • | |

- **UNIX FreeBSD scheduling**

| Priority Class | Thread Type | Description |
|--|--------------------|---|
| 0–63 | Bottom-half kernel | Scheduled by interrupts. Can block to await a resource |
| 64–127 | Top-half kernel | Runs until blocked or done. Can block to await a resource |
| 128–159 | Real-time user | Allowed to run until blocked or until a higher-priority thread becomes available. Preemptive scheduling |
| 160–223 | Time-sharing user | Adjusts priorities based on processor usage |
| 224–255 | Idle user | Only run when there are no time sharing or real-time threads to run |
| Note: Lower number corresponds to higher priority. | | |