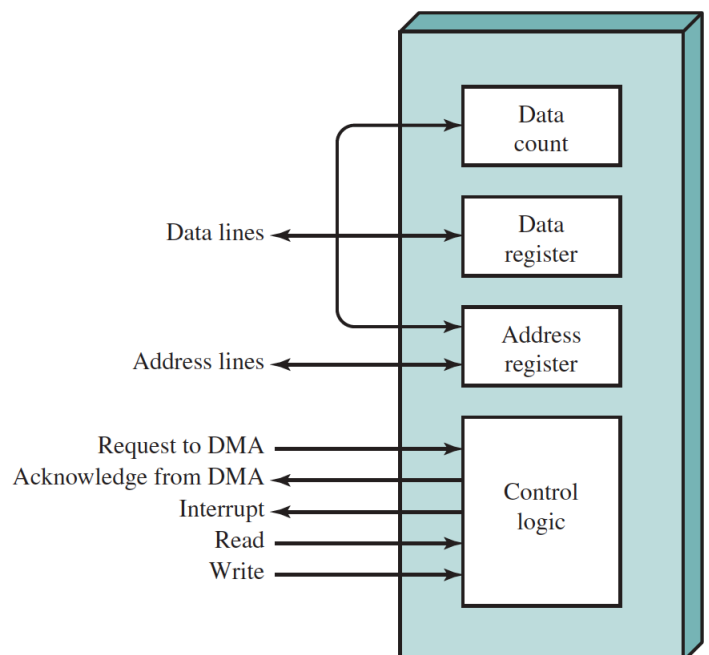


Chapter 11




- Categories of I/O devices:
 - Human readable (keyboard, printer)
 - Machine readable (sensors, disk drives)
 - Communication (modems)
- Properties of I/O devices:
 - Data rate
 - Application – the use affects its software
 - Complexity of control
 - Unit of transfer
 - Data representation
 - Error conditions
- Techniques of performing I/O:

	No Interrupts	Use of Interrupts
I/O-to-Memory Transfer through Processor	Programmed I/O	Interrupt-driven I/O
Direct I/O-to-Memory Transfer		Direct memory access (DMA)

- Evolution of the I/O function:
 - Processor directly controls peripheral device
 - → programmed I/O module
 - → I/O module w/ interrupts
 - → direct memory access
 - → I/O channel, can execute programs on its own
 - → I/O processor, has its own memory
- **Direct memory access**
 - Processor sends to DMA:
 - Read or write
 - Address of I/O device
 - Starting address
 - # of words
 - DMA transfers data directly between I/O and memory
 - DMA sends interrupt to processor
 - Having DMA share the system bus w/ the processor is inefficient
 - DMA can be directly connected with one or more I/O devices
 - DMA can share I/O bus with all devices



- OS design issues:
 - Efficiency** – I/O is slow; tend to be bottlenecks
 - Generality** – desirable to handle all devices in a uniform manner
- Layers of I/O function:
 - Logical I/O – deals with the device as a resource/model
 - Interacts via commands, e.g. open, close, read, write
 - Device I/O – converts operations & data into I/O commands
 - Scheduling & control – queueing & scheduling of I/O operations; interrupts
 - Interacts with actual device hardware
 - In a peripheral device:
 - User process → logical I/O → device I/O → scheduling & control → hardware*
 - Directory management – converts symbolic file names to identifiers
 - File system – logical organization of files; user operations, e.g. open, close, read, write
 - Physical organization – converts logical file references to physical storage addresses
 - In a file system:
 - User process → directory mgmt. → file system → physical org. → device I/O → scheduling & control → hardware*
- Block-oriented device – data transfers are made one block at a time
- Stream-oriented device – data is transferred in streams of bytes
- I/O buffering**
 - Pages of memory involved in I/O need to be locked in MM during I/O
 - E.g. I/O transfer → user memory; user process cannot be swapped out
 - Single buffer – OS assigns a buffer in system memory
 - Block-oriented
 - Input transfer → buffer; buffer block → user space; request another block
 - OS can process one block while next block is being read in
 - Swapping can occur since transfer is in system memory (not user)
 - Stream-oriented
 - Buffer reads in/writes out one line at a time
 - Double buffer
 - Process transfers data to/from one buffer while the OS empties/fills the other buffer
 - Circular buffer
 - Process & OS cycle between >2 buffers




- **Disk scheduling**
- Disk performance
 - Seek time – time taken to position the head of the track
 - Rotational delay – time for beginning of sector to reach the head
 - Access time = seek time + rotational delay
 - Transfer time – time taken to transfer data
 - Sequential access is *much faster* than random access
- Disk scheduling policies
 - FIFO – access disk in the same order the requests were received (fair)
 - Priority – not intended to optimize but to meet OS objectives
 - Starvation possible
 - Last in first out – takes advantage of locality, improves throughput
 - Starvation possible
 - Shortest service time first – select request that requires the least disk arm movement (minimum seek time)
 - Starvation possible
 - SCAN – arm moves in one direction and satisfies all requests along the way; reverse when finished
 - Biased against area most recently visited
 - Circular SCAN – scans in one directory, arm returns to opposite end after finishing
 - N-step SCAN – segments disk request queue into sub-queues of length N
 - Process sub-queues using SCAN
 - $N = 1 \rightarrow$ FIFO
- **RAID** – redundant array of independent disks
 - Set of physical disks viewed by OS as a single logical drive
 - Data are distributed across the disks by striping
 - Redundant disk capacity stores parity information, ensures data recoverability


Category	Level	Description	Disks Required	Data Availability	Large I/O Data Transfer Capacity	Small I/O Request Rate
Striping	0	Nonredundant	N	Lower than single disk	Very high	Very high for both read and write
Mirroring	1	Mirrored	$2N$	Higher than RAID 2, 3, 4, or 5; lower than RAID 6	Higher than single disk for read; similar to single disk for write	Up to twice that of a single disk for read; similar to single disk for write
Parallel access	2	Redundant via Hamming code	$N + m$	Much higher than single disk; comparable to RAID 3, 4, or 5	Highest of all listed alternatives	Approximately twice that of a single disk
	3	Bit-interleaved parity	$N + 1$	Much higher than single disk; comparable to RAID 2, 4, or 5	Highest of all listed alternatives	Approximately twice that of a single disk
Independent access	4	Block-interleaved parity	$N + 1$	Much higher than single disk; comparable to RAID 2, 3, or 5	Similar to RAID 0 for read; significantly lower than single disk for write	Similar to RAID 0 for read; significantly lower than single disk for write
	5	Block-interleaved distributed parity	$N + 1$	Much higher than single disk; comparable to RAID 2, 3, or 4	Similar to RAID 0 for read; lower than single disk for write	Similar to RAID 0 for read; generally lower than single disk for write
	6	Block-interleaved dual distributed parity	$N + 2$	Highest of all listed alternatives	Similar to RAID 0 for read; lower than RAID 5 for write	Similar to RAID 0 for read; significantly lower than RAID 5 for write

Note: N , number of data disks; m , proportional to $\log N$.


- Striping: (strip 0, 1, 2, 3 form a stripe)



(a) RAID 0 (nonredundant)




(b) RAID 1 (mirrored)




(c) RAID 2 (redundancy through Hamming code)


- Redundancy by parity – parity disk stores bits from which lost data can be recovered




(d) RAID 3 (bit-interleaved parity)



(e) RAID 4 (block-level parity)



(f) RAID 5 (block-level distributed parity)

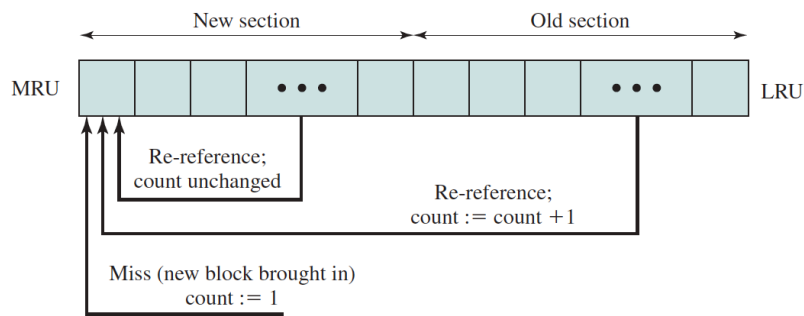


(g) RAID 6 (dual redundancy)


• Disk cache

▪ Replacement policy

- Least recently used
- Least frequently used – block with fewest references
- Frequency based replacement



(a) FIFO



(b) Use of three sections

• UNIX SVR4 I/O

- Buffered I/O – uses buffer cache (disk cache)
- Unbuffered I/O – uses DMA

• Linux I/O

- Disk scheduling – uses the elevator scheduler
 - Deadline scheduler
 - Anticipatory I/O scheduler

• Windows I/O

- Synchronous vs. asynchronous I/O

Chapter 12

-