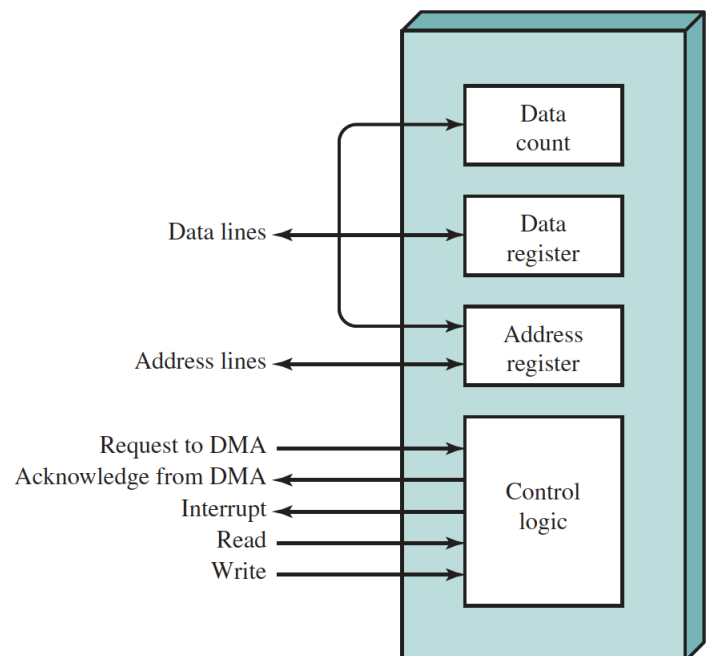


Chapter 11 (11.1 – 11.9)

- Categories of I/O devices:
 - Human readable (keyboard, printer)
 - Machine readable (sensors, disk drives)
 - Communication (modems)
- Properties of I/O devices:
 - Data rate
 - Application – the use affects its software
 - Complexity of control
 - Unit of transfer
 - Data representation
 - Error conditions
- Techniques of performing I/O:

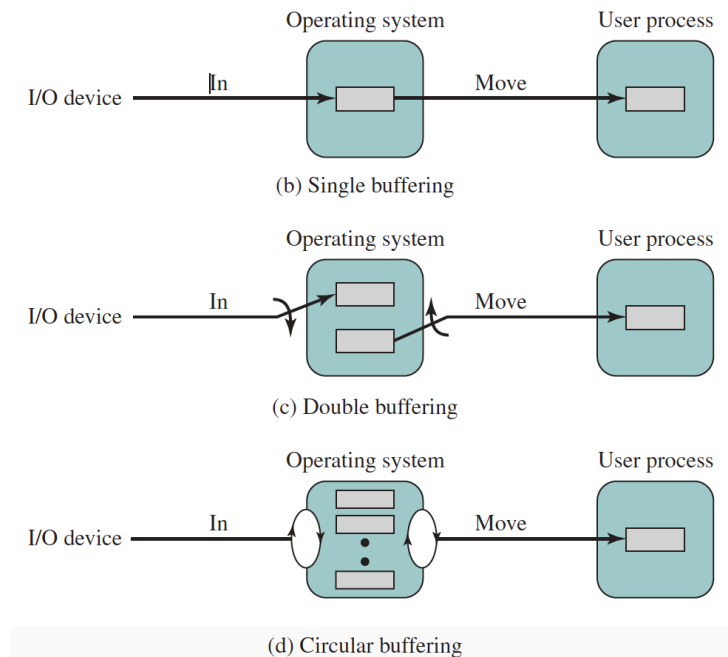
| | No Interrupts | Use of Interrupts |
|--|----------------|----------------------------|
| I/O-to-Memory Transfer through Processor | Programmed I/O | Interrupt-driven I/O |
| Direct I/O-to-Memory Transfer | | Direct memory access (DMA) |

- **Direct memory access**
 - Processor sends to DMA:
 - Read or write
 - Address of I/O device
 - Starting address
 - # of words (length)
 - DMA transfers data directly between I/O and memory
 - DMA sends interrupt to processor
 - Having DMA share the system bus w/ the processor is inefficient
 - DMA can be directly connected with one or more I/O devices
 - DMA can share I/O bus with all devices



- OS design issues:
 - Efficiency – I/O is slow; tend to be bottlenecks
 - Generality – desirable to handle all devices in a uniform manner
- Layers of I/O function:
 - Logical I/O – deals with the device as a resource/model
 - Interacts via commands, e.g. open, close, read, write
 - Device I/O – converts operations & data into I/O commands
 - Scheduling & control – queueing & scheduling of I/O operations; interrupts

- Interacts with actual device hardware
- In a peripheral device:
 - User process → logical I/O → device I/O → scheduling & control → hardware*
- Directory management – converts symbolic file names to identifiers
- File system – logical organization of files; user operations, e.g. open, close, read, write
- Physical organization – converts logical file references to physical storage addresses
- In a file system:
 - User process → directory mgmt. → file system → physical org. → device I/O → scheduling & control → hardware*
- Block-oriented device – data transfers are made one block at a time
- Stream-oriented device – data is transferred in streams of bytes
- I/O buffering**
 - Pages of memory involved in I/O need to be locked in MM during I/O
 - E.g. I/O transfer → user memory; user process cannot be swapped out
 - Solution = buffering
 - Disadvantage: deadlock is possible
 - Single buffer – OS assigns a buffer in system memory
 - Block-oriented
 - Input transfer → buffer; buffer block → user space; request another block
 - OS can process one block while next block is being read in
 - Swapping can occur since transfer is in system memory (not user)
 - Stream-oriented
 - Buffer reads in/writes out one line at a time
 - Double buffer
 - Process transfers data to/from one buffer while the OS empties/fills the other buffer
 - Circular buffer
 - Process & OS cycle between >2 buffers
 - Good for short bursts of I/O



- Disk scheduling**
- Disk performance
 - Seek time – time taken to position the head of the track
 - Rotational delay – time for beginning of sector to reach the head
 - Access time = seek time + rotational delay
 - Transfer time – time taken to transfer data

- Sequential access is *much faster* than random access
- Disk scheduling policies

| Policy | Description/advantages | Disadvantages |
|------------------------------------|--|--|
| FIFO | <ul style="list-style-type: none"> • Treats requests fairly | <ul style="list-style-type: none"> • Performance approaches random access if # of processes is high |
| Priority | <ul style="list-style-type: none"> • Not intended to optimize but to meet OS objectives | <ul style="list-style-type: none"> • Starvation possible |
| LIFO | <ul style="list-style-type: none"> • Takes advantage of locality → high throughput | <ul style="list-style-type: none"> • Starvation possible |
| Shortest service time first (SSTF) | <ul style="list-style-type: none"> • Minimizes disk arm movement/seek time | <ul style="list-style-type: none"> • Starvation possible |
| SCAN/elevator | <ul style="list-style-type: none"> • Arm moves in one direction and satisfies all requests along the way; reverse when finished • Deterministic | <ul style="list-style-type: none"> • Biased against area most recently visited; arm stickiness |
| LOOK | <ul style="list-style-type: none"> • Arm reverses at the last request (not the end of track) | |
| Circular-SCAN | <ul style="list-style-type: none"> • Only scan in only direction → reduces delay for new requests • Deterministic | <ul style="list-style-type: none"> • Arm stickiness |
| N-step SCAN | <ul style="list-style-type: none"> • Segment request queue into sub-queues of length N • Sub-queues use SCAN; main queue is FIFO • Reduces arm stickiness | |
| FSCAN | <ul style="list-style-type: none"> • Two queues for old & new requests • New requests are deferred until old requests are serviced | |

- Performance:
 - $SSTF \approx SCAN < C-SCAN < FIFO$
- **RAID** – redundant array of independent disks
 - Set of physical disks viewed by OS as a single logical drive
 - Data are distributed across the disks by striping
 - Logically contiguous strips are spread over multiple disks → can be accessed in parallel, reducing transfer time
 - Redundant disk capacity stores parity information, ensures data recoverability
 - Redundancy by parity – parity disk stores bits from which lost data can be recovered


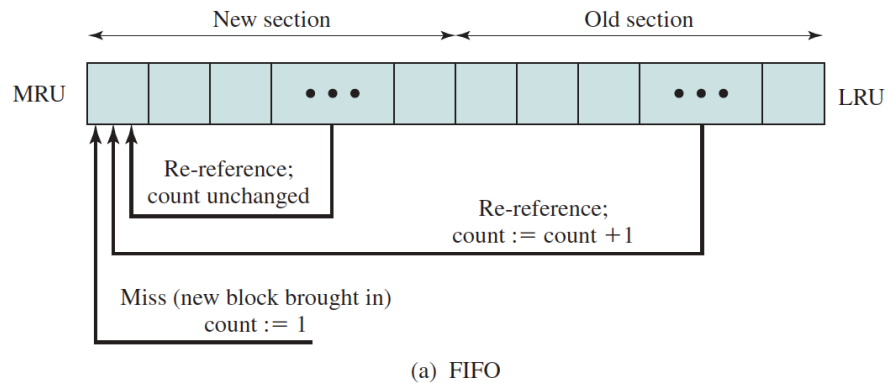
| Category | Level | Description | Disks Required | Data Availability | Large I/O Data Transfer Capacity | Small I/O Request Rate |
|--------------------|-------|---|----------------|---|--|--|
| Striping | 0 | Nonredundant | N | Lower than single disk | Very high | Very high for both read and write |
| Mirroring | 1 | Mirrored | $2N$ | Higher than RAID 2, 3, 4, or 5; lower than RAID 6 | Higher than single disk for read; similar to single disk for write | Up to twice that of a single disk for read; similar to single disk for write |
| Parallel access | 2 | Redundant via Hamming code | $N + m$ | Much higher than single disk; comparable to RAID 3, 4, or 5 | Highest of all listed alternatives | Approximately twice that of a single disk |
| | 3 | Bit-interleaved parity | $N + 1$ | Much higher than single disk; comparable to RAID 2, 4, or 5 | Highest of all listed alternatives | Approximately twice that of a single disk |
| Independent access | 4 | Block-interleaved parity | $N + 1$ | Much higher than single disk; comparable to RAID 2, 3, or 5 | Similar to RAID 0 for read; significantly lower than single disk for write | Similar to RAID 0 for read; significantly lower than single disk for write |
| | 5 | Block-interleaved distributed parity | $N + 1$ | Much higher than single disk; comparable to RAID 2, 3, or 4 | Similar to RAID 0 for read; lower than single disk for write | Similar to RAID 0 for read; generally lower than single disk for write |
| | 6 | Block-interleaved dual distributed parity | $N + 2$ | Highest of all listed alternatives | Similar to RAID 0 for read; lower than RAID 5 for write | Similar to RAID 0 for read; significantly lower than RAID 5 for write |

Note: N , number of data disks; m , proportional to $\log N$.

• Disk cache

▪ Replacement policy

- LRU
- Least frequently used – block with fewest references
- Frequency based replacement:



• UNIX SVR4 I/O

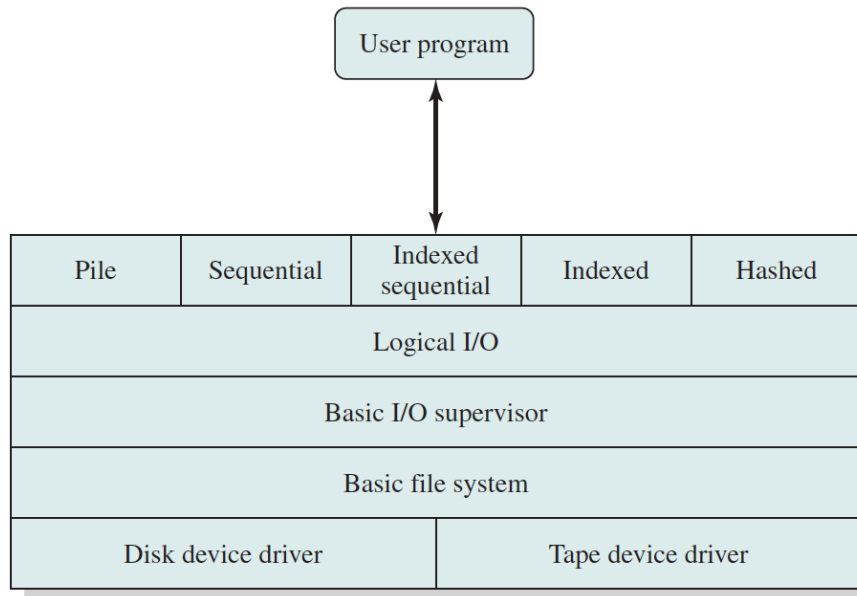
- Buffered I/O – uses buffer cache (disk cache)
- Unbuffered I/O – uses DMA

• Linux I/O

- Disk scheduling – uses the elevator scheduler (variation of LOOK)
 - Deadline scheduler
 - Anticipatory I/O scheduler

Chapter 12 (12.1 – 12.2, 12.4 – 12.8)

- Properties of files:
 - Long-term existence
 - Sharable between processes
 - Structure
- File structure:
 - Field → record → file → database
- **File system architecture**



- **File organization**
 - Criteria:
 - Short access time
 - Ease of update
 - Economy of storage
 - Simple maintenance
 - Reliability
 - Pile file
 - Chronological list of variable-length records
 - Record access is exhaustive (linear)
 - Sequential file
 - Fixed-length records with fixed fields (i.e. table)
 - Sorted by a key field
 - Record access is exhaustive
 - Indexed sequential file
 - Index – pointer into file that allows faster random access
 - Only key field is indexed
 - Indexed file
 - Exhaustive index – every record is indexed

- Partial index – points to area of interest
- Hashed file
 - No ordering; every key is hashed

| File Method | Space Attributes | | Update Record Size | | Retrieval | | |
|--------------------|------------------|-------|--------------------|---------|---------------|--------|------------|
| | Variable | Fixed | Equal | Greater | Single record | Subset | Exhaustive |
| Pile | A | B | A | E | E | D | B |
| Sequential | F | A | D | F | F | D | A |
| Indexed sequential | F | B | B | D | B | D | B |
| Indexed | B | C | C | C | A | B | D |
| Hashed | F | B | B | F | B | F | E |

A = Excellent, well suited to this purpose $\approx O(r)$
 B = Good $\approx O(o \times r)$
 C = Adequate $\approx O(r \log n)$
 D = Requires some extra effort $\approx O(n)$
 E = Possible with extreme effort $\approx O(r \times n)$
 F = Not reasonable for this purpose $\approx O(n^2)$


where

r = size of the result
 o = number of records that overflow
 n = number of records in file

• Record blocking



Fixed blocking



Variable blocking: spanned



Variable blocking: unspanned

- File allocation

| | Contiguous | Chained | Indexed | |
|----------------------------------|------------|--------------|--------------|----------|
| Preallocation? | Necessary | Possible | Possible | |
| Fixed or Variable Size Portions? | Variable | Fixed blocks | Fixed blocks | Variable |
| Portion Size | Large | Small | Small | Medium |
| Allocation Frequency | Once | Low to high | High | Low |
| Time to Allocate | Medium | Long | Short | Medium |
| File Allocation Table Size | One entry | One entry | Large | Medium |

-